

# [1] Amazon Fine Food Reviews Analysis

Data Source: <https://www.kaggle.com/snap/amazon-fine-food-reviews> (<https://www.kaggle.com/snap/amazon-fine-food-reviews>)

The Amazon Fine Food Reviews dataset consists of reviews of fine foods from Amazon.

Number of reviews: 568,454

Number of users: 256,059

Number of products: 74,258

Timespan: Oct 1999 - Oct 2012

Number of Attributes/Columns in data: 10

Attribute Information:

1. Id
2. ProductId - unique identifier for the product
3. UserId - unique identifier for the user
4. ProfileName
5. HelpfulnessNumerator - number of users who found the review helpful
6. HelpfulnessDenominator - number of users who indicated whether they found the review helpful or not
7. Score - rating between 1 and 5
8. Time - timestamp for the review
9. Summary - brief summary of the review
10. Text - text of the review

## Objective:

Given a review, determine whether the review is positive (Rating of 4 or 5) or negative (rating of 1 or 2).

[Q] How to determine if a review is positive or negative?

[Ans] We could use the Score/Rating. A rating of 4 or 5 could be considered a positive review. A review of 1 or 2 could be considered negative. A review of 3 is neutral and ignored. This is an approximate and proxy way of determining the polarity (positivity/negativity) of a review.

## [7.1] Loading the data

The dataset is available in two forms

1. .csv file
2. SQLite Database

In order to load the data, We have used the SQLITE dataset as it is easier to query the data and visualise the data efficiently.

Here as we only want to get the global sentiment of the recommendations (positive or negative), we will purposefully ignore all Scores equal to 3. If the score is above 3, then the recommendation will be set to "positive". Otherwise, it will be set to "negative".

In [1]:

```

import warnings
from sklearn.exceptions import DataConversionWarning
warnings.filterwarnings(action='ignore', category=DataConversionWarning)
warnings.filterwarnings(action='ignore', category=UserWarning)
warnings.filterwarnings(action='ignore', category=FutureWarning)

import math
import random
import traceback
import sqlite3
import itertools
import datetime as dt
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sn
from tqdm import tqdm
from sklearn import preprocessing
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split
from prettytable import PrettyTable
from sklearn.metrics import accuracy_score, precision_score, recall_score, confusion_matrix, classification_report
from sklearn.metrics import make_scorer
from scipy.stats import uniform
from scipy.sparse import find
from sklearn.externals import joblib
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import RandomizedSearchCV
from sklearn.model_selection import TimeSeriesSplit
from gensim.models import Word2Vec
from gensim.models import KeyedVectors
from gensim.models import word2vec

```

```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\externals\joblib\__init___.py:15: DeprecationWarning: sklearn.externals.joblib is deprecated in 0.21 and will be removed in 0.23. Please import this functionality directly from joblib, which can be installed with: pip install joblib. If this warning is raised when loading pickled models, you may need to re-serialize those models with scikit-learn 0.21+.

```

```
warnings.warn(msg, category=DeprecationWarning)
```

## Exploratory Data Analysis

### [7.1.2] Data Cleaning: Deduplication

It is observed (as shown in the table below) that the reviews data had many duplicate entries. Hence it was necessary to remove duplicates in order to get unbiased results for the analysis of the data. Following is an example:

In [2]:

```
#Using sqlite3 to retrieve data from sqlite file

con = sqlite3.connect("final.sqlite")#Loading Cleaned/ Preprocesed text that we did in Text

#Using pandas functions to query from sql table
final = pd.read_sql_query("""
SELECT * FROM Reviews order by Time
""",con)

#Reviews is the name of the table given
#Taking only the data where score != 3 as score 3 will be neutral and it won't help us much
final.head()
```

Out[2]:

	index	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	Helpfu
0	138706	150524	0006641040	ACITT7DI6IDDL	shari zychinski	0	
1	138683	150501	0006641040	AJ46FKXOVC7NR	Nicholas A Mesiano	2	
2	417839	451856	B00004CXX9	AIUWLEQ1ADEG5	Elizabeth Medina	0	
3	346055	374359	B00004CI84	A344SMIA5JECGM	Vincent P. Ross	1	
4	417838	451855	B00004CXX9	AJH6LUC1UT1ON	The Phantom of the Opera	0	

In [3]:

```
#Sorting data according to ProductId in ascending order
sorted_data=final.sort_values('ProductId', axis=0, ascending=True, inplace=False, kind='quicksort')

#Deduplication of entries
final=sorted_data.drop_duplicates(subset={"UserId","ProfileName","Time","Text"}, keep='first')
```

In [4]:

```
final = final[final.HelpfulnessNumerator <= final.HelpfulnessDenominator]
```

## Text Preprocessing: Stemming, stop-word removal and Lemmatization

In [5]:

```
import nltk
#set of stopwords in English
from nltk.corpus import stopwords
stop = set(stopwords.words('english'))
words_to_keep = set(('not'))
stop -= words_to_keep
#initialising the snowball stemmer
sno = nltk.stem.SnowballStemmer('english')

#function to clean the word of any html-tags
def cleanhtml(sentence):
    cleanr = re.compile('<.*?>')
    cleantext = re.sub(cleanr, ' ', sentence)
    return cleantext

#function to clean the word of any punctuation or special characters
def cleanpunc(sentence):
    cleaned = re.sub(r'[?|!|\'|\"|#]', r'', sentence)
    cleaned = re.sub(r'[\.,|)|(|\|/]', r' ', cleaned)
    return cleaned
```

In [6]:

```
import re
#Code for removing HTML tags , punctuations . Code for removing stopwords . Code for checking
# also greater than 2 . Code for stemming and also to convert them to lowercase letters
i=0
str1=' '
final_string=[]
all_positive_words=[] # store words from +ve reviews here
all_negative_words=[] # store words from -ve reviews here.
s=''
for sent in final['Text'].values:
    filtered_sentence=[]
    #print(sent);
    sent=cleanhtml(sent) # remove HTML tags
    for w in sent.split():
        for cleaned_words in cleanpunc(w).split():
            if((cleaned_words.isalpha()) & (len(cleaned_words)>2)):
                if(cleaned_words.lower() not in stop):
                    s=(sno.stem(cleaned_words.lower())).encode('utf8')
                    filtered_sentence.append(s)
                    if (final['Score'].values)[i] == 'positive':
                        all_positive_words.append(s) #list of all words used to describe positive
                    if(final['Score'].values)[i] == 'negative':
                        all_negative_words.append(s) #list of all words used to describe negative
                else:
                    continue
            else:
                continue

    str1 = b" ".join(filtered_sentence) #final string of cleaned words

    final_string.append(str1)
    i+=1
```

In [7]:

```
final['CleanedText']=final_string  
final['CleanedText']=final['CleanedText'].str.decode("utf-8")  
#below the processed review can be seen in the CleanedText Column  
final.head()
```

Out[7]:

	index	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator
<b>288755</b>	7137	7802	B00474H936	A2JT9TJ74UYK30	Laura	2
<b>20630</b>	479208	518195	B001E5E3LW	A156JLBSURBJE3	R. Corcoran "barryvillage2"	4
<b>190985</b>	210138	227746	B0033HPPIO	A3FLZCOLQMZ6J1	MrsKaraT	0
<b>125128</b>	140196	152146	B0002EOW4E	A1NE4STC8I0T6D	Jana Dengler	0
<b>142392</b>	9789	10707	B00389Q4XW	A23CDAI9ZJWVKO	Elaine	0



In [8]:

```

n_samples = 25000
final = final.sample(n_samples)

###Sorting as we want according to time series
final.sort_values('Time',inplace=True)
final.head(10)

```

Out[8]:

	index	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator
0	138706	150524	0006641040	ACITT7DI6IDDL	shari zychinski	0	
9	346094	374400	B00004CI84	A2DEE7F9XKP3ZR	jerome	0	
21	138018	149790	B00004S1C6	A1IU7S4HCK1XK0	Joanna Daneman	25	
24	246412	274410	B00004CI84	A2D0BL1B4KBO	"neradice found"	2	

In [9]:

```

def plot_confusion_matrix(test_y, predict_y):
    C = confusion_matrix(test_y, predict_y)
    A = (((C.T)/(C.sum(axis=1))).T)
    B = (C/C.sum(axis=0))
    plt.figure(figsize=(20,4))
    labels = [0,1]
    # representing A in heatmap format
    cmap=sns.light_palette("blue")
    plt.subplot(1, 3, 1)
    sns.heatmap(C, annot=True, cmap=cmap, fmt=".3f", xticklabels=labels, yticklabels=labels)
    plt.xlabel('Predicted Class')
    plt.ylabel('Original Class')
    plt.title("Confusion matrix")

    plt.subplot(1, 3, 2)
    sns.heatmap(B, annot=True, cmap=cmap, fmt=".3f", xticklabels=labels, yticklabels=labels)
    plt.xlabel('Predicted Class')
    plt.ylabel('Original Class')
    plt.title("Precision matrix")

    plt.subplot(1, 3, 3)
    # representing B in heatmap format
    sns.heatmap(A, annot=True, cmap=cmap, fmt=".3f", xticklabels=labels, yticklabels=labels)
    plt.xlabel('Predicted Class')
    plt.ylabel('Original Class')
    plt.title("Recall matrix")

    plt.show()

```

# Bag of Words (BoW)

In [10]:

```
x = final['CleanedText'].values
y = final['Score']

# split the data set into train and test
X_train, X_test, Y_train, Y_test = train_test_split(x, y, test_size=0.3, random_state=0)
```

In [11]:

```
count_vect = CountVectorizer(min_df = 10)
X_train_vec = count_vect.fit_transform(X_train)
X_test_vec = count_vect.transform(X_test)
print("the type of count vectorizer :",type(X_train_vec))
print("the shape of out text BOW vectorizer : ",X_train_vec.get_shape())
print("the number of unique words :", X_train_vec.get_shape()[1])
```

```
the type of count vectorizer : <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text BOW vectorizer : (17500, 3790)
the number of unique words : 3790
```

## Random Forest(RF) with GridSearchCV

In [12]:

```

from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import accuracy_score, confusion_matrix, f1_score, precision_score, recall

base_learners = [5,10,15,20,25,50,75,100,125,150,175,200,225,250,275,300]

depth = [5,10,20,30,40,50,60,70,80,90,100,110,120,130,140,150]
param_grid = {'n_estimators': base_learners, 'max_depth': depth}
RFC = RandomForestClassifier(max_features='sqrt')
model = GridSearchCV(RFC, param_grid, scoring = 'f1_weighted', cv=3 , n_jobs = -1, pre_dispatch='all')
model.fit(X_train_vec, Y_train)
print("Model with best parameters :\n", model.best_estimator_)
print("Accuracy of the model : ", model.score(X_test_vec, Y_test))
# Optimal value of number of base learners
optimal_learners = model.best_estimator_.n_estimators
print("The optimal number of base learners is : ", optimal_learners)

optimal_depth = model.best_estimator_.max_depth
print("The optimal number of depth is : ", optimal_depth)

```

Model with best parameters :

```

RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                        max_depth=120, max_features='sqrt', max_leaf_nodes=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=1, min_samples_split=2,
                        min_weight_fraction_leaf=0.0, n_estimators=10, n_jobs=None,
                        oob_score=False, random_state=None, verbose=0,
                        warm_start=False)

```

Accuracy of the model : 0.854847789068937

The optimal number of base learners is : 10

The optimal number of depth is : 120

In [13]:

```
import seaborn as sns
```



In [14]:


```

print("Best HyperParameter: ",model.best_params_)
print(model.best_score_)
scores = model.cv_results_['mean_test_score'].reshape(len(base_learners),len(depth))

plt.figure(figsize=(16, 12))
sns.heatmap(scores, annot=True, cmap=plt.cm.hot, fmt=".3f", xticklabels=base_learners, yticklabels=depth)
plt.xlabel('n_estimators')
plt.ylabel('max_depth')
plt.xticks(np.arange(len(base_learners)), base_learners)
plt.yticks(np.arange(len(depth)), depth)
plt.title('Grid Search f1 Score')
plt.show()

```

Best HyperParameter: {'max\_depth': 120, 'n\_estimators': 10}  
 0.8522059515926641



In [15]:


```
import seaborn as sns
```

In [16]:

```
#confusion matrix,precision matrix,recall matrix,accuracy
from sklearn.metrics import accuracy_score, precision_recall_fscore_support, f1_score
rf = RandomForestClassifier(n_estimators=optimal_learners, max_features='sqrt', n_jobs=-1)
rf.fit(X_train_vec,Y_train)
Y_pred = rf.predict(X_test_vec)
Y_test_accuracy = accuracy_score(Y_test, Y_pred, normalize=True, sample_weight=None)*100
print('Accuracy of the model at optimal hyperparameter depth = %d is:  %f%%' % (optimal_learners, Y_test_accuracy))
print('Confusion matrix for the model is:')
plot_confusion_matrix(Y_test, Y_pred)
f1score= f1_score(Y_test, Y_pred, average='weighted')
print('f1 score value for the model is: %s'% f1score)
precisionscore=precision_score(Y_test, Y_pred,pos_label='positive' )
print('precision score for the model is: %s'% precisionscore)
y_train_pred = rf.predict(X_train_vec)
Y_train_accuracy =accuracy_score(Y_train, y_train_pred, normalize=True, sample_weight=None)
plot_confusion_matrix(Y_train, y_train_pred)
print('Accuracy of the model at optimal hyperparameter depth = %d is:  %f%%' % (optimal_learners, Y_train_accuracy))
f1score= f1_score(Y_train, y_train_pred, average='weighted')
print('f1 score value for the model is: %s'% f1score)
precisionscore=precision_score(Y_train, y_train_pred,pos_label='positive' )
print('precision score for the model is: %s'% precisionscore)
```


Accuracy of the model at optimal hyperparameter depth = 10 is: 87.773333%

Confusion matrix for the model is:



f1 score value for the model is: 0.8631446618220829

precision score for the model is: 0.8965970498028334



Accuracy of the model at optimal hyperparameter depth = 10 is: 99.525714%

f1 score value for the model is: 0.9952355106965103

precision score for the model is: 0.995136449608214

## Gradient Boosting Decision Tree(GBDT) with GridSearchCV

In [17]:

```

# Importing library
from sklearn.ensemble import GradientBoostingClassifier

base_learners = [5,10,15,20,25,50,75,100,125,150,175,200,225,250,275,300]
depth = [5,10,20,30,40,50,60,70,80,90,100,110,120,130,140,150]

param_grid = {'n_estimators': base_learners, 'max_depth': depth}
GBC = GradientBoostingClassifier(max_features='sqrt',subsample=0.1)
model = GridSearchCV(GBC, param_grid, scoring = 'f1_weighted', cv=3, n_jobs = -1,pre_dispatch=5)
model.fit(X_train_vec, Y_train)
print("Model with best parameters :\n",model.best_estimator_)
print("Accuracy of the model : ",model.score(X_test_vec, Y_test))

# Optimal value of number of base learners
optimal_learners = model.best_estimator_.n_estimators
print("The optimal number of base learners is : ",optimal_learners)

# Optimal value of Learning rate
optimal_depth = model.best_estimator_.max_depth
print("\nThe optimal value of max depth is : ",optimal_depth)

```

Model with best parameters :

```

GradientBoostingClassifier(criterion='friedman_mse', init=None,
                           learning_rate=0.1, loss='deviance', max_depth=20,
                           max_features='sqrt', max_leaf_nodes=None,
                           min_impurity_decrease=0.0, min_impurity_split=None,
                           min_samples_leaf=1, min_samples_split=2,
                           min_weight_fraction_leaf=0.0, n_estimators=300,
                           n_iter_no_change=None, presort='auto', random_state=None,
                           subsample=0.1, tol=0.0001, validation_fraction=0.1,
                           verbose=0, warm_start=False)

```

Accuracy of the model : 0.8733301309514601

The optimal number of base learners is : 300

The optimal value of max depth is : 20

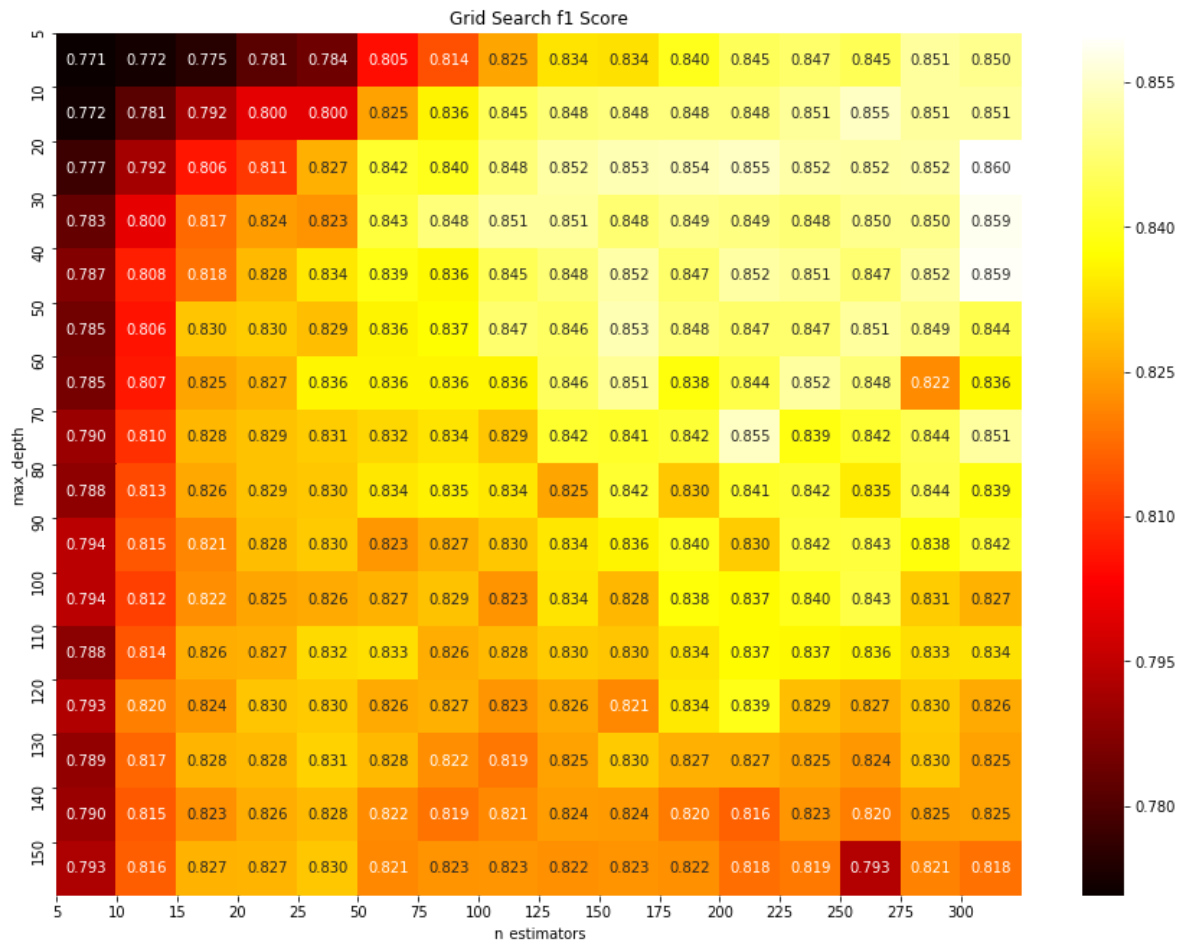
In [18]:

```

scores = model.cv_results_['mean_test_score'].reshape(len(base_learners),len(depth))

plt.figure(figsize=(15, 11))
sns.heatmap(scores, annot=True, cmap=plt.cm.hot, fmt=".3f", xticklabels=base_learners, yticklabels=depth)
plt.xlabel('n_estimators')
plt.ylabel('max_depth')
plt.xticks(np.arange(len(base_learners)), base_learners)
plt.yticks(np.arange(len(depth)), depth)
plt.title('Grid Search f1 Score')
plt.show()

```





In [19]:


```

import seaborn as sns
#confusion matrix,precision matrix,recall matrix,accuracy
from sklearn.metrics import accuracy_score, precision_recall_fscore_support, f1_score
gb = GradientBoostingClassifier(n_estimators=optimal_learners, max_depth=optimal_depth, max
gb.fit(X_train_vec,Y_train)
Y_pred = gb.predict(X_test_vec)
Y_test_accuracy = accuracy_score(Y_test, Y_pred, normalize=True, sample_weight=None)*100
print('Accuracy of the model at optimal hyperparameter estimator = %d is:  %f%%' % (optimal
print('Confusion matrix for the model is:')
plot_confusion_matrix(Y_test, Y_pred)
f1score= f1_score(Y_test, Y_pred, average='weighted')
print('f1 score value for  the model is: %s'% f1score)
precisionscore=precision_score(Y_test, Y_pred,pos_label='positive' )
print('precision score for  the model is: %s'% precisionscore)
y_train_pred = gb.predict(X_train_vec)
Y_train_accuracy =accuracy_score(Y_train, y_train_pred, normalize=True, sample_weight=None)
plot_confusion_matrix(Y_train, y_train_pred)
print('Accuracy of the model at optimal hyperparameter estimator = %d is:  %f%%' % (optimal
f1score= f1_score(Y_train, y_train_pred, average='weighted')
print('f1 score value for  the model is: %s'% f1score)
precisionscore=precision_score(Y_train, y_train_pred,pos_label='positive' )
print('precision score for  the model is: %s'% precisionscore)

```


Accuracy of the model at optimal hyperparameter estimator = 300 is: 86.960000%

Confusion matrix for the model is:



f1 score value for the model is: 0.8635895384846582

precision score for the model is: 0.9083029197080292



Accuracy of the model at optimal hyperparameter estimator = 300 is: 92.611429%

f1 score value for the model is: 0.9232447769192669

precision score for the model is: 0.9420270803207572

# Term Frequency - Inverse Document Frequency (TF-IDF)

In [20]:

```
tf_idf_vect = TfidfVectorizer(min_df=10)
X_train_vec = tf_idf_vect.fit_transform(X_train)
X_test_vec = tf_idf_vect.transform(X_test)
print("the type of count vectorizer :",type(X_train_vec))
print("the shape of out text TFIDF vectorizer : ",X_train_vec.get_shape())
print("the number of unique words :", X_train_vec.get_shape()[1])
```

```
the type of count vectorizer : <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text TFIDF vectorizer : (17500, 3790)
the number of unique words : 3790
```

## Random Forest(RF) with GridSearchCV

In [21]:

```
base_learners = [5,10,15,20,25,50,75,100,125,150,175,200,225,250,275,300]

depth = [5,10,20,30,40,50,60,70,80,90,100,110,120,130,140,150]
param_grid = {'n_estimators': base_learners, 'max_depth': depth}
RFC = RandomForestClassifier(max_features='sqrt')
model = GridSearchCV(RFC, param_grid, scoring = 'f1_weighted', cv=3 , n_jobs = -1,pre_dispatch=5)
model.fit(X_train_vec, Y_train)
print("Model with best parameters :\n",model.best_estimator_)
print("Accuracy of the model : ",model.score(X_test_vec, Y_test))

# Optimal value of number of base learners
optimal_learners = model.best_estimator_.n_estimators
print("The optimal number of base learners is : ",optimal_learners)

optimal_depth=model.best_estimator_.max_depth
print("The optimal number of depth is : ",optimal_depth)
```

Model with best parameters :

```
RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                        max_depth=130, max_features='sqrt', max_leaf_nodes=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=1, min_samples_split=2,
                        min_weight_fraction_leaf=0.0, n_estimators=10, n_jobs=None,
                        oob_score=False, random_state=None, verbose=0,
                        warm_start=False)
```

Accuracy of the model : 0.8614202901162389

The optimal number of base learners is : 10

The optimal number of depth is : 130



In [22]:


```

print("Best HyperParameter: ",model.best_params_)
print(model.best_score_)
scores = model.cv_results_['mean_test_score'].reshape(len(base_learners),len(depth))

plt.figure(figsize=(16, 12))
sns.heatmap(scores, annot=True, cmap=plt.cm.hot, fmt=".3f", xticklabels=base_learners, yticklabels=depth)
plt.xlabel('n_estimators')
plt.ylabel('max_depth')
plt.xticks(np.arange(len(base_learners)), base_learners)
plt.yticks(np.arange(len(depth)), depth)
plt.title('Grid Search f1 Score')
plt.show()

```


Best HyperParameter: {'max\_depth': 130, 'n\_estimators': 10}  
 0.8483462681980961




In [23]:

```
#confusion matrix,precision matrix,recall matrix,accuracy
from sklearn.metrics import accuracy_score, precision_recall_fscore_support, f1_score
rf = RandomForestClassifier(n_estimators=optimal_learners,max_depth=optimal_depth, max_feat
rf.fit(X_train_vec,Y_train)
Y_pred = rf.predict(X_test_vec)
Y_test_accuracy = accuracy_score(Y_test, Y_pred, normalize=True, sample_weight=None)*100
print('Accuracy of the model at optimal hyperparameter depth = %d is: %f%%' % (optimal_lea
print('Confusion matrix for the model is:')
plot_confusion_matrix(Y_test, Y_pred)
f1score= f1_score(Y_test, Y_pred, average='weighted')
print('f1 score value for the model is: %s'% f1score)
precisionscore=precision_score(Y_test, Y_pred,pos_label='positive' )
print('precision score for the model is: %s'% precisionscore)
y_train_pred = rf.predict(X_train_vec)
Y_train_accuracy =accuracy_score(Y_train, y_train_pred, normalize=True, sample_weight=None)
plot_confusion_matrix(Y_train, y_train_pred)
print('Accuracy of the model at optimal hyperparameter depth = %d is: %f%%' % (optimal_lea
f1score= f1_score(Y_train, y_train_pred, average='weighted')
print('f1 score value for the model is: %s'% f1score)
precisionscore=precision_score(Y_train, y_train_pred,pos_label='positive' )
print('precision score for the model is: %s'% precisionscore)
```

Accuracy of the model at optimal hyperparameter depth = 10 is: 87.946667%  
 Confusion matrix for the model is:



f1 score value for the model is: 0.8578714371127336  
 precision score for the model is: 0.8879703534777651



Accuracy of the model at optimal hyperparameter depth = 10 is: 98.422857%  
 f1 score value for the model is: 0.9838967456985855  
 precision score for the model is: 0.9818799546998868

## Gradient Boosting Decision Tree(GBDT) with GridSearchCV

In [24]:

```
base_learners = [5,10,15,20,25,50,75,100,125,150,175,200,225,250,275,300]
depth = [5,10,20,30,40,50,60,70,80,90,100,110,120,130,140,150]

param_grid = {'n_estimators': base_learners, 'max_depth': depth}
GBC = GradientBoostingClassifier(max_features='sqrt', subsample=0.1)
model = GridSearchCV(GBC, param_grid, scoring = 'f1_weighted', cv=3 , n_jobs = -1, pre_dispatch='all')
model.fit(X_train_vec, Y_train)
print("Model with best parameters :\n", model.best_estimator_)
print("Accuracy of the model : ", model.score(X_test_vec, Y_test))

# Optimal value of number of base learners
optimal_learners = model.best_estimator_.n_estimators
print("The optimal number of base learners is : ", optimal_learners)

# Optimal value of depth
optimal_depth = model.best_estimator_.max_depth
print("\nThe optimal value of depth is : ", optimal_depth)
```

Model with best parameters :

```
GradientBoostingClassifier(criterion='friedman_mse', init=None,
                           learning_rate=0.1, loss='deviance', max_depth=10,
                           max_features='sqrt', max_leaf_nodes=None,
                           min_impurity_decrease=0.0, min_impurity_split=None,
                           min_samples_leaf=1, min_samples_split=2,
                           min_weight_fraction_leaf=0.0, n_estimators=275,
                           n_iter_no_change=None, presort='auto', random_state=None,
                           subsample=0.1, tol=0.0001, validation_fraction=0.1,
                           verbose=0, warm_start=False)
```

Accuracy of the model : 0.868906442646433

The optimal number of base learners is : 275

The optimal value of depth is : 10

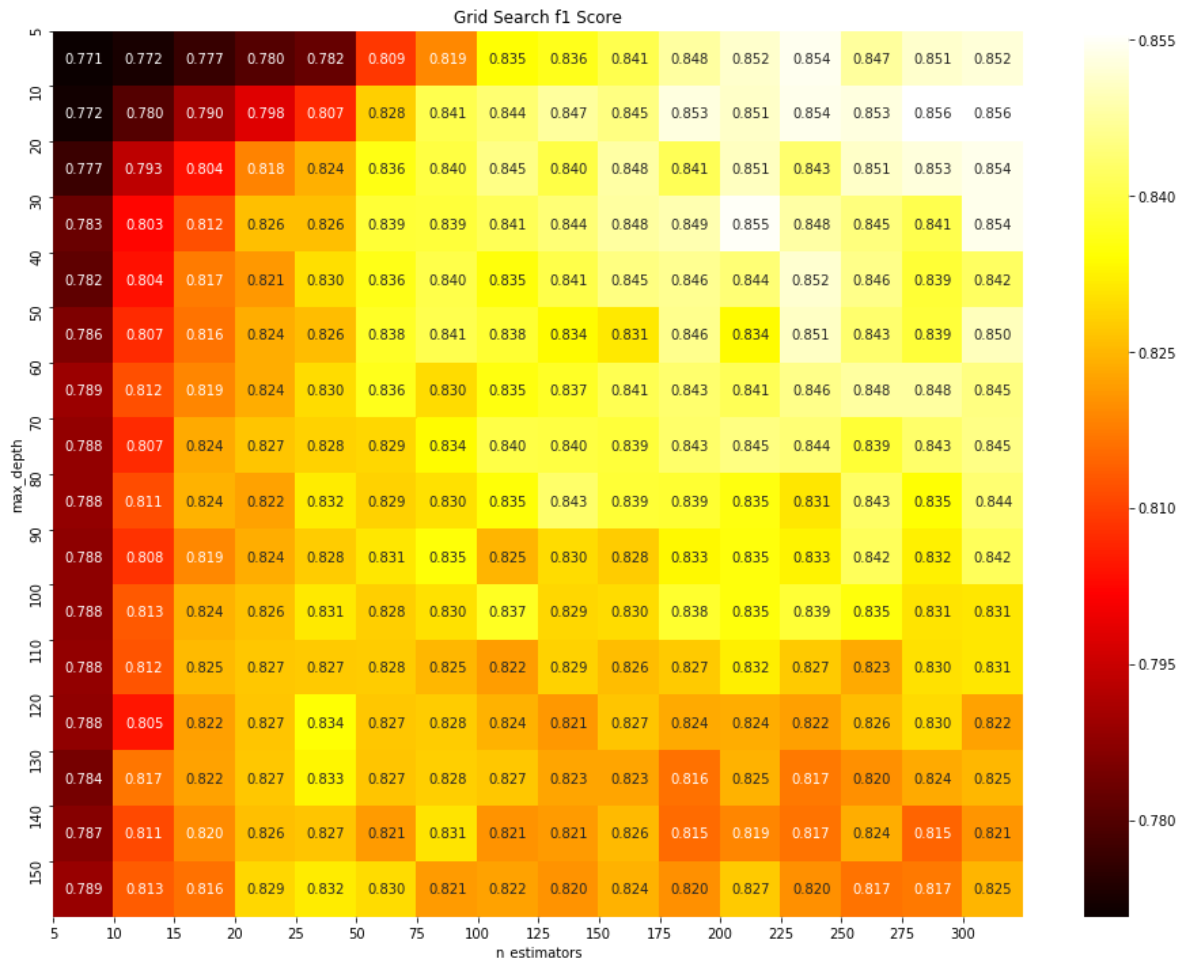
In [25]:

```

scores = model.cv_results_['mean_test_score'].reshape(len(base_learners),len(depth))

plt.figure(figsize=(16, 12))
sns.heatmap(scores, annot=True, cmap=plt.cm.hot, fmt=".3f", xticklabels=base_learners, yticklabels=depth)
plt.xlabel('n_estimators')
plt.ylabel('max_depth')
plt.xticks(np.arange(len(base_learners)), base_learners)
plt.yticks(np.arange(len(depth)), depth)
plt.title('Grid Search f1 Score')
plt.show()

```




In [26]:

```
#confusion matrix,precision matrix,recall matrix,accuracy
from sklearn.metrics import accuracy_score, precision_recall_fscore_support, f1_score
gb = GradientBoostingClassifier(n_estimators=optimal_learners, max_depth=optimal_depth, max
gb.fit(X_train_vec,Y_train)
Y_pred = gb.predict(X_test_vec)
Y_test_accuracy = accuracy_score(Y_test, Y_pred, normalize=True, sample_weight=None)*100
print('Accuracy of the model at optimal hyperparameter estimator = %d is: %f%%' % (optimal
print('Confusion matrix for the model is:')
plot_confusion_matrix(Y_test, Y_pred)
f1score= f1_score(Y_test, Y_pred, average='weighted')
print('f1 score value for the model is: %s'% f1score)
precisionscore=precision_score(Y_test, Y_pred,pos_label='positive' )
print('precision score for the model is: %s'% precisionscore)
y_train_pred = gb.predict(X_train_vec)
Y_train_accuracy =accuracy_score(Y_train, y_train_pred, normalize=True, sample_weight=None)
plot_confusion_matrix(Y_train, y_train_pred)
print('Accuracy of the model at optimal hyperparameter estimator = %d is: %f%%' % (optimal
f1score= f1_score(Y_train, y_train_pred, average='weighted')
print('f1 score value for the model is: %s'% f1score)
precisionscore=precision_score(Y_train, y_train_pred,pos_label='positive' )
print('precision score for the model is: %s'% precisionscore)
```


Accuracy of the model at optimal hyperparameter estimator = 275 is: 88.280000%

Confusion matrix for the model is:



f1 score value for the model is: 0.8708263999703252

precision score for the model is: 0.9026652922986306



Accuracy of the model at optimal hyperparameter estimator = 275 is: 91.160000%

f1 score value for the model is: 0.9043376888246839

precision score for the model is: 0.9219209617598159

# Average Word2Vec

In [27]:

```
# List of sentence in X_train text
sent_of_train=[]
for sent in X_train:
    sent_of_train.append(sent.split())

# List of sentence in X_test text
sent_of_test=[]
for sent in X_test:
    sent_of_test.append(sent.split())

# Train your own Word2Vec model using your own train text corpus
# min_count = 5 considers only words that occurred atleast 5 times
w2v_model=Word2Vec(sent_of_train,min_count=5,size=50, workers=4)

w2v_words = list(w2v_model.wv.vocab)
print("number of words that occurred minimum 5 times ",len(w2v_words))
```

number of words that occurred minimum 5 times 5870

In [28]:

```
# compute average word2vec for each review for X_train .
train_vectors = [];
for sent in sent_of_train:
    sent_vec = np.zeros(50)
    cnt_words =0;
    for word in sent: #
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    train_vectors.append(sent_vec)
```

In [29]:

```
# compute average word2vec for each review for X_test .
test_vectors = [];
for sent in sent_of_test:
    sent_vec = np.zeros(50)
    cnt_words =0;
    for word in sent: #
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    test_vectors.append(sent_vec)

X_train_vec = train_vectors
X_test_vec = test_vectors
```

# Random Forest(RF) with GridSearchCV

In [30]:

```
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import accuracy_score, confusion_matrix, f1_score, precision_score, recall
base_learners = [5,10,15,20,25,50,75,100,125,150,175,200,225,250,275,300]

depth = [5,10,20,30,40,50,60,70,80,90,100,110,120,130,140,150]
param_grid = {'n_estimators': base_learners, 'max_depth': depth}
RFC = RandomForestClassifier(max_features=0.2)
model = GridSearchCV(RFC, param_grid, scoring = 'f1_micro', cv=3 , n_jobs = -1, pre_dispatch='all')
model.fit(X_train_vec, Y_train)
print("Model with best parameters :\n", model.best_estimator_)
print("Accuracy of the model : ", model.score(X_test_vec, Y_test))

# Optimal value of number of base learners
optimal_learners = model.best_estimator_.n_estimators
print("The optimal number of base learners is : ", optimal_learners)

optimal_depth = model.best_estimator_.max_depth
print("The optimal number of depth is : ", optimal_depth)
```

Model with best parameters :

```
RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                        max_depth=80, max_features=0.2, max_leaf_nodes=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=1, min_samples_split=2,
                        min_weight_fraction_leaf=0.0, n_estimators=200, n_jobs=None,
                        oob_score=False, random_state=None, verbose=0,
                        warm_start=False)
```

Accuracy of the model : 0.8716000000000002

The optimal number of base learners is : 200

The optimal number of depth is : 80

In [32]:


```

print("Best HyperParameter: ", model.best_params_)
print(model.best_score_)
scores = model.cv_results_['mean_test_score'].reshape(len(base_learners), len(depth))

plt.figure(figsize=(16, 12))
sns.heatmap(scores, annot=True, cmap=plt.cm.hot, fmt=".3f", xticklabels=base_learners, yticklabels=depth)
plt.xlabel('n_estimators')
plt.ylabel('max_depth')
plt.xticks(np.arange(len(base_learners)), base_learners)
plt.yticks(np.arange(len(depth)), depth)
plt.title('Grid Search f1 Score')
plt.show()

```

Best HyperParameter: {'max\_depth': 80, 'n\_estimators': 200}  
 0.8686857142857143






In [33]:

```
#confusion matrix,precision matrix,recall matrix,accuracy
from sklearn.metrics import accuracy_score, precision_recall_fscore_support, f1_score
rf = RandomForestClassifier(n_estimators=optimal_learners,max_depth=optimal_depth, max_feat
rf.fit(X_train_vec,Y_train)
Y_pred = rf.predict(X_test_vec)
Y_test_accuracy = accuracy_score(Y_test, Y_pred, normalize=True, sample_weight=None)*100
print('Accuracy of the model at optimal hyperparameter depth = %d is: %f%%' % (optimal_lea
print('Confusion matrix for the model is:')
plot_confusion_matrix(Y_test, Y_pred)
f1score= f1_score(Y_test, Y_pred, average='micro')
print('f1 score value for the model is: %s'% f1score)
precisionscore=precision_score(Y_test, Y_pred,pos_label='positive' )
print('precision score for the model is: %s'% precisionscore)
y_train_pred = rf.predict(X_train_vec)
Y_train_accuracy =accuracy_score(Y_train, y_train_pred, normalize=True, sample_weight=None)
plot_confusion_matrix(Y_train, y_train_pred)
print('Accuracy of the model at optimal hyperparameter depth = %d is: %f%%' % (optimal_lea
f1score= f1_score(Y_train, y_train_pred, average='micro')
print('f1 score value for the model is: %s'% f1score)
precisionscore=precision_score(Y_train, y_train_pred,pos_label='positive' )
print('precision score for the model is: %s'% precisionscore)
```


Accuracy of the model at optimal hyperparameter depth = 200 is: 86.84000  
0%

Confusion matrix for the model is:



f1 score value for the model is: 0.8684

precision score for the model is: 0.8753675955748494



## Gradient Boosting Decision Tree(GBDT) with GridSearchCV

In [34]:

```

from sklearn.ensemble import GradientBoostingClassifier
base_learners = [5,10,15,20,25,50,75,100,125,150,175,200,225,250,275,300]
depth = [5,10,20,30,40,50,60,70,80,90,100,110,120,130,140,150]

param_grid = {'n_estimators': base_learners, 'max_depth': depth}
GBC = GradientBoostingClassifier(max_features='sqrt',subsample=0.1)
model = GridSearchCV(GBC, param_grid, scoring = 'f1_weighted', cv=3 , n_jobs = -1,pre_dispatch=5)
model.fit(X_train_vec, Y_train)
print("Model with best parameters :\n",model.best_estimator_)
print("Accuracy of the model : ",model.score(X_test_vec, Y_test))

# Optimal value of number of base learners
optimal_learners = model.best_estimator_.n_estimators
print("The optimal number of base learners is : ",optimal_learners)

# Optimal value of depth
optimal_depth = model.best_estimator_.max_depth
print("\nThe optimal value of depth is : ",optimal_depth)

```

Model with best parameters :

```

GradientBoostingClassifier(criterion='friedman_mse', init=None,
                           learning_rate=0.1, loss='deviance', max_depth=5,
                           max_features='sqrt', max_leaf_nodes=None,
                           min_impurity_decrease=0.0, min_impurity_split=None,
                           min_samples_leaf=1, min_samples_split=2,
                           min_weight_fraction_leaf=0.0, n_estimators=75,
                           n_iter_no_change=None, presort='auto', random_state=None,
                           subsample=0.1, tol=0.0001, validation_fraction=0.1,
                           verbose=0, warm_start=False)

```

Accuracy of the model : 0.8457429290695962

The optimal number of base learners is : 75

The optimal value of depth is : 5


In [35]:

```

scores = model.cv_results_['mean_test_score'].reshape(len(base_learners),len(depth))

plt.figure(figsize=(16, 12))
sns.heatmap(scores, annot=True, cmap=plt.cm.hot, fmt=".3f", xticklabels=base_learners, yticklabels=depth)
plt.xlabel('n_estimators')
plt.ylabel('max_depth')
plt.xticks(np.arange(len(base_learners)), base_learners)
plt.yticks(np.arange(len(depth)), depth)
plt.title('Grid Search f1 Score')
plt.show()

```




In [36]:


```
#confusion matrix,precision matrix,recall matrix,accuracy
from sklearn.metrics import accuracy_score, precision_recall_fscore_support, f1_score
gb = GradientBoostingClassifier(n_estimators=optimal_learners, max_depth=optimal_depth, max
gb.fit(X_train_vec,Y_train)
Y_pred = gb.predict(X_test_vec)
Y_test_accuracy = accuracy_score(Y_test, Y_pred, normalize=True, sample_weight=None)*100
print('Accuracy of the model at optimal hyperparameter estimator = %d is:  %f%%' % (optimal
print('Confusion matrix for the model is:')
plot_confusion_matrix(Y_test, Y_pred)
f1score= f1_score(Y_test, Y_pred, average='weighted')
print('f1 score value for the model is: %s'% f1score)
precisionscore=precision_score(Y_test, Y_pred,pos_label='positive' )
print('precision score for the model is: %s'% precisionscore)
y_train_pred = gb.predict(X_train_vec)
Y_train_accuracy =accuracy_score(Y_train, y_train_pred, normalize=True, sample_weight=None)
plot_confusion_matrix(Y_train, y_train_pred)
print('Accuracy of the model at optimal hyperparameter estimator = %d is:  %f%%' % (optimal
f1score= f1_score(Y_train, y_train_pred, average='weighted')
print('f1 score value for the model is: %s'% f1score)
precisionscore=precision_score(Y_train, y_train_pred,pos_label='positive' )
print('precision score for the model is: %s'% precisionscore)
```

Accuracy of the model at optimal hyperparameter estimator = 75 is: 86.693333%

Confusion matrix for the model is:



f1 score value for the model is: 0.8484711895572471  
precision score for the model is: 0.8873443382565884



Accuracy of the model at optimal hyperparameter estimator = 75 is: 87.977143%

f1 score value for the model is: 0.8654642542025487  
precision score for the model is: 0.8964803312629399

# TFIDF-Word2Vec

In [37]:

```
# TF-IDF weighted Word2Vec
tf_idf_vect = TfidfVectorizer()

# final_tf_idf1 is the sparse matrix with row= sentence, col=word and cell_val = tfidf
final_tf_idf1 = tf_idf_vect.fit_transform(X_train)

# tfidf words/col-names
tfidf_feat = tf_idf_vect.get_feature_names()

# compute TFIDF Weighted Word2Vec for each review for X_test .
tfidf_test_vectors = [];
row=0;
for sent in sent_of_test:
    sent_vec = np.zeros(50)
    weight_sum =0;
    for word in sent:
        if word in w2v_words:
            vec = w2v_model.wv[word]
            # obtain the tf_idfidf of a word in a sentence/review
            tf_idf = final_tf_idf1[row, tfidf_feat.index(word)]
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    tfidf_test_vectors.append(sent_vec)
    row += 1
```

In [38]:

```
# compute TFIDF Weighted Word2Vec for each review for X_train .
tfidf_train_vectors = [];
row=0;
for sent in sent_of_train:
    sent_vec = np.zeros(50)
    weight_sum =0;
    for word in sent:
        if word in w2v_words:
            vec = w2v_model.wv[word]
            # obtain the tf_idfidf of a word in a sentence/review
            tf_idf = final_tf_idf1[row, tfidf_feat.index(word)]
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    tfidf_train_vectors.append(sent_vec)
    row += 1

X_train_vec = tfidf_train_vectors
X_test_vec = tfidf_test_vectors
```

## Random Forest(RF) with GridSearchCV

In [39]:

```
base_learners = [5,10,15,20,25,50,75,100,125,150,175,200,225,250,275,300]

depth = [5,10,20,30,40,50,60,70,80,90,100,110,120,130,140,150]
param_grid = {'n_estimators': base_learners, 'max_depth': depth}
RFC = RandomForestClassifier(max_features='sqrt')
model = GridSearchCV(RFC, param_grid, scoring = 'f1_micro', cv=3 , n_jobs = -1,pre_dispatch
model.fit(X_train_vec, Y_train)
print("Model with best parameters :\n",model.best_estimator_)
print("Accuracy of the model : ",model.score(X_test_vec, Y_test))

# Optimal value of number of base learners
optimal_learners = model.best_estimator_.n_estimators
print("The optimal number of base learners is : ",optimal_learners)

optimal_depth=model.best_estimator_.max_depth
print("The optimal number of depth is : ",optimal_depth)
```

Model with best parameters :

```
RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
    max_depth=130, max_features='sqrt', max_leaf_nodes=None,
    min_impurity_decrease=0.0, min_impurity_split=None,
    min_samples_leaf=1, min_samples_split=2,
    min_weight_fraction_leaf=0.0, n_estimators=225, n_jobs=None,
    oob_score=False, random_state=None, verbose=0,
    warm_start=False)
```

Accuracy of the model : 0.8050666666666666

The optimal number of base learners is : 225

The optimal number of depth is : 130

In [40]:


```

print("Best HyperParameter: ",model.best_params_)
print(model.best_score_)
scores = model.cv_results_['mean_test_score'].reshape(len(base_learners),len(depth))

plt.figure(figsize=(16, 12))
sns.heatmap(scores, annot=True, cmap=plt.cm.hot, fmt=".3f", xticklabels=base_learners, yticklabels=depth)
plt.xlabel('n_estimators')
plt.ylabel('max_depth')
plt.xticks(np.arange(len(base_learners)), base_learners)
plt.yticks(np.arange(len(depth)), depth)
plt.title('Grid Search f1 Score')
plt.show()

```


Best HyperParameter: {'max\_depth': 130, 'n\_estimators': 225}  
 0.8602857142857143




In [41]:

```
#confusion matrix,precision matrix,recall matrix,accuracy
from sklearn.metrics import accuracy_score, precision_recall_fscore_support, f1_score
rf = RandomForestClassifier(n_estimators=optimal_learners,max_depth=optimal_depth, max_feat
rf.fit(X_train_vec,Y_train)
Y_pred = rf.predict(X_test_vec)
Y_test_accuracy = accuracy_score(Y_test, Y_pred, normalize=True, sample_weight=None)*100
print('Accuracy of the model at optimal hyperparameter depth = %d is: %f%%' % (optimal_lea
print('Confusion matrix for the model is:')
plot_confusion_matrix(Y_test, Y_pred)
f1score= f1_score(Y_test, Y_pred, average='micro')
print('f1 score value for the model is: %s'% f1score)
precisionscore=precision_score(Y_test, Y_pred,pos_label='positive' )
print('precision score for the model is: %s'% precisionscore)
y_train_pred = rf.predict(X_train_vec)
Y_train_accuracy =accuracy_score(Y_train, y_train_pred, normalize=True, sample_weight=None)
plot_confusion_matrix(Y_train, y_train_pred)
print('Accuracy of the model at optimal hyperparameter depth = %d is: %f%%' % (optimal_lea
f1score= f1_score(Y_train, y_train_pred, average='micro')
print('f1 score value for the model is: %s'% f1score)
precisionscore=precision_score(Y_train, y_train_pred,pos_label='positive' )
print('precision score for the model is: %s'% precisionscore)
```

Accuracy of the model at optimal hyperparameter depth = 225 is: 81.026667%  
 Confusion matrix for the model is:



f1 score value for the model is: 0.8102666666666667  
 precision score for the model is: 0.8553903882234088



Accuracy of the model at optimal hyperparameter depth = 225 is: 100.000000%  
 f1 score value for the model is: 1.0  
 precision score for the model is: 1.0

## Gradient Boosting Decision Tree(GBDT) with



# GridSearchCV

In [42]:

```
base_learners = [5,10,15,20,25,50,75,100,125,150,175,200,225,250,275,300]
depth = [5,10,20,30,40,50,60,70,80,90,100,110,120,130,140,150]

param_grid = {'n_estimators': base_learners, 'max_depth': depth}
GBC = GradientBoostingClassifier(max_features='sqrt',subsample=0.1)
model = GridSearchCV(GBC, param_grid, scoring = 'f1_micro', cv=3 , n_jobs = -1,pre_dispatch
model.fit(X_train_vec, Y_train)
print("Model with best parameters :\n",model.best_estimator_)
print("Accuracy of the model : ",model.score(X_test_vec, Y_test))

# Optimal value of number of base learners
optimal_learners = model.best_estimator_.n_estimators
print("The optimal number of base learners is : ",optimal_learners)

# Optimal value of depth
optimal_depth = model.best_estimator_.max_depth
print("\nThe optimal value of depth is : ",optimal_depth)
```

Model with best parameters :

```
GradientBoostingClassifier(criterion='friedman_mse', init=None,
    learning_rate=0.1, loss='deviance', max_depth=5,
    max_features='sqrt', max_leaf_nodes=None,
    min_impurity_decrease=0.0, min_impurity_split=None,
    min_samples_leaf=1, min_samples_split=2,
    min_weight_fraction_leaf=0.0, n_estimators=25,
    n_iter_no_change=None, presort='auto', random_state=None,
    subsample=0.1, tol=0.0001, validation_fraction=0.1,
    verbose=0, warm_start=False)
```

Accuracy of the model : 0.7290666666666666

The optimal number of base learners is : 25

The optimal value of depth is : 5

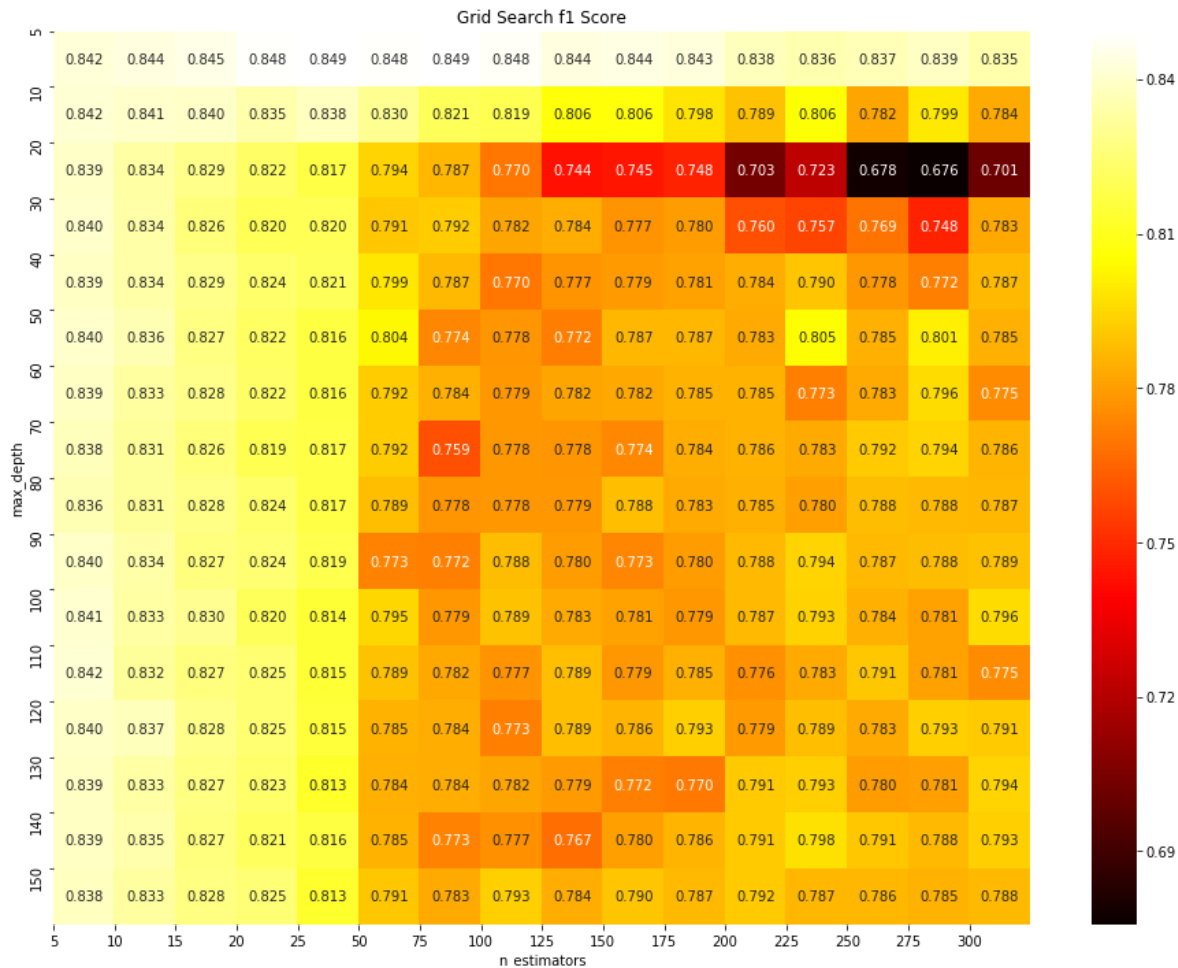
In [43]:

```

scores = model.cv_results_['mean_test_score'].reshape(len(base_learners),len(depth))

plt.figure(figsize=(16, 12))
sns.heatmap(scores, annot=True, cmap=plt.cm.hot, fmt=".3f", xticklabels=base_learners, yticklabels=depth)
plt.xlabel('n_estimators')
plt.ylabel('max_depth')
plt.xticks(np.arange(len(base_learners)), base_learners)
plt.yticks(np.arange(len(depth)), depth)
plt.title('Grid Search f1 Score')
plt.show()

```




In [44]:

```
from sklearn.metrics import accuracy_score, precision_recall_fscore_support, f1_score
gb = GradientBoostingClassifier(n_estimators=optimal_learners, max_depth=optimal_depth, max
gb.fit(X_train_vec,Y_train)
Y_pred = gb.predict(X_test_vec)
Y_test_accuracy = accuracy_score(Y_test, Y_pred, normalize=True, sample_weight=None)*100
print('Accuracy of the model at optimal hyperparameter estimator = %d is: %f%%' % (optimal
print('Confusion matrix for the model is:')
plot_confusion_matrix(Y_test, Y_pred)
f1score= f1_score(Y_test, Y_pred, average='micro')
print('f1 score value for the model is: %s'% f1score)
precisionscore=precision_score(Y_test, Y_pred,pos_label='positive' )
print('precision score for the model is: %s'% precisionscore)
y_train_pred = gb.predict(X_train_vec)
Y_train_accuracy =accuracy_score(Y_train, y_train_pred, normalize=True, sample_weight=None)
plot_confusion_matrix(Y_train, y_train_pred)
print('Accuracy of the model at optimal hyperparameter estimator = %d is: %f%%' % (optimal
f1score= f1_score(Y_train, y_train_pred, average='micro')
print('f1 score value for the model is: %s'% f1score)
precisionscore=precision_score(Y_train, y_train_pred,pos_label='positive' )
print('precision score for the model is: %s'% precisionscore)
```


Accuracy of the model at optimal hyperparameter estimator = 25 is: 73.266667%

Confusion matrix for the model is:



f1 score value for the model is: 0.7326666666666668

precision score for the model is: 0.8616153205661948



Accuracy of the model at optimal hyperparameter estimator = 25 is: 85.548571%

f1 score value for the model is: 0.8554857142857143

precision score for the model is: 0.8614892861370901

In [49]:

```

from prettytable import PrettyTable

# Names of models
featurization = ['Bag of Words', 'Bag of Words', 'TFIDF ', 'TFIDF ', 'avg w2v', 'avg w2vw', 'TFIDF
model=['Random Forest(RF) ', 'Gradient Boosting Decision Tree', 'Random Forest(RF) ', 'Gradi

baselearners =[10,300,10,275,200,75,225,25]
depth=[120,20,130,10,80,5,130,5]
accuracy = [87.77,86.96,87.94,88.28,86.84,86.69,81.02,73.26]
F1score= [0.8631,0.8635,0.8578,0.8708,0.8684,0.8484,0.8102,0.7326]
precision=[0.8965,0.9083,0.8879,0.9026,0.8753,0.8873,0.8553,0.8616]
numbering = [1,2,3,4,5,6,7,8]

# Initializing prettytable
ptable = PrettyTable()

# Adding columns
ptable.add_column("S.NO.", numbering)
ptable.add_column("MODEL", featurization)
ptable.add_column("Method", model)
ptable.add_column("base learners", baselearners)
ptable.add_column("depth", depth)
ptable.add_column("accuracy", accuracy)
ptable.add_column("f1score", F1score)
ptable.add_column("precision", precision)

# Printing the Table
print(ptable)

```

S.NO.	MODEL	Method	base learners	depth	accuracy	f1score	precision
1	Bag of Words	Random Forest(RF)	10	120	87.77	0.8631	0.8965
2	Bag of Words	Gradient Boosting Decision Tree	300	20	86.96	0.8635	0.9083
3	TFIDF	Random Forest(RF)	10	130	87.94	0.8578	0.8879
4	TFIDF	Gradient Boosting Decision Tree	275	10	88.28	0.8708	0.9026
5	avg w2v	Random Forest(RF)	200	80	86.84	0.8684	0.8753
6	avg w2vw	Gradient Boosting Decision Tree	75	5	86.69	0.8484	0.8873
7	TFIDF weighted word2vec	Random Forest(RF)	25	130	81.02	0.8102	0.8553
8	TFIDF weighted word2vec	Gradient Boosting Decision Tree	25	5	73.26	0.7326	0.8616

In [ ]:

In [ ]:

In [ ]: